Prompting Audios Using Acoustic Properties for Emotion Representation

Hira Dhamyal, Benjamin Elizalde, Soham Deshmukh, Huaming Wang, Bhiksha Raj, Rita Singh

Introduction

- Emotions lie on a continuum, but current models treat emotions as a finite valued discrete variable. This representation does not capture the diversity in the expression of emotion. **To better represent them, we propose the use of natural language descriptions (or prompts)**.
- In this work, we address the challenge **of automatically generating these prompts** and training a model to better learn emotion representations from **audio and prompt pairs**.
- We use **acoustic properties** that are correlated to emotion like pitch, intensity, speech rate, and articulation rate to automatically generate prompts, i.e., '**acoustic prompts**'.

DATASETS

Dataset	# Files	# Classes	Emotions
CMU-MOSEI	23K	9	ang, exc, fear, sad, frus, neu, sur, hap, dis
IEMOCAP	10K	9	hap, fear, sad, sur, exc, ang, neu, disappoint, frus
MELD	10K	7	neu, sur, fear, sad, joy, disgust, ang
CREMA-D	7K	6	ang, dis, fear, hap, neu, sad
RAVDESS	2.5K	8	neu, calm, hap, sad, ang, fear, disgust, sur
CMU-MOSI	2.2K	3	neu, positive, negative

MOTIVATION

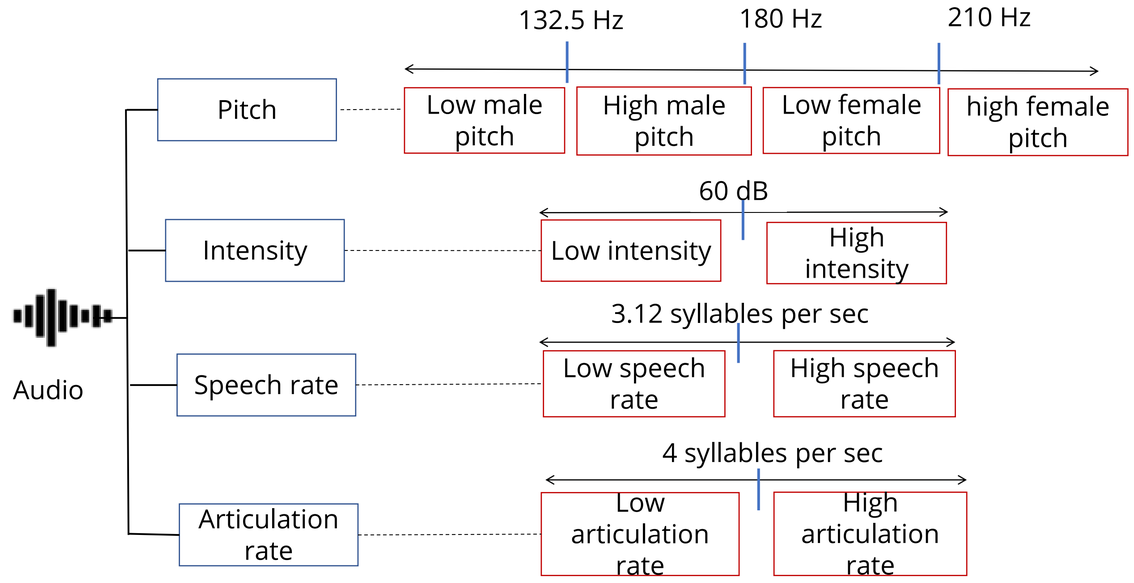

METHODOLOGY

EXAMPLE PROMPTS

Property	Prompt
Class Label	• {emotion}
Pitch	• High female pitch {emotion} • High male pitch {emotion}
Intensity	• High intensity {emotion}
Speech Rate	• High speech rate {emotion}
Articulation Rate	• High articulation rate {emotion}

MODEL

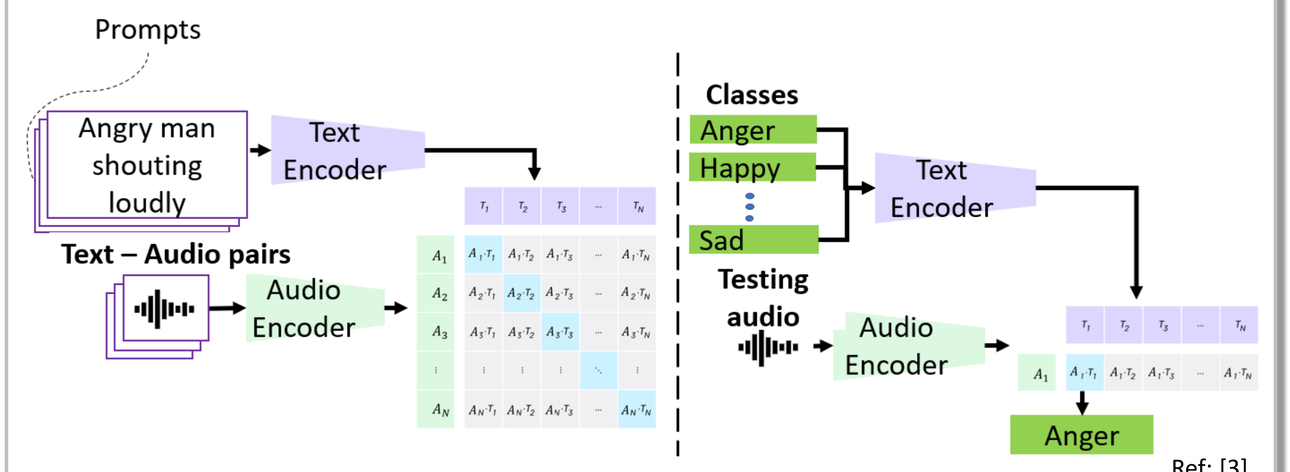

Ref: [3]

RESULTS – Emotion Audio Retrieval

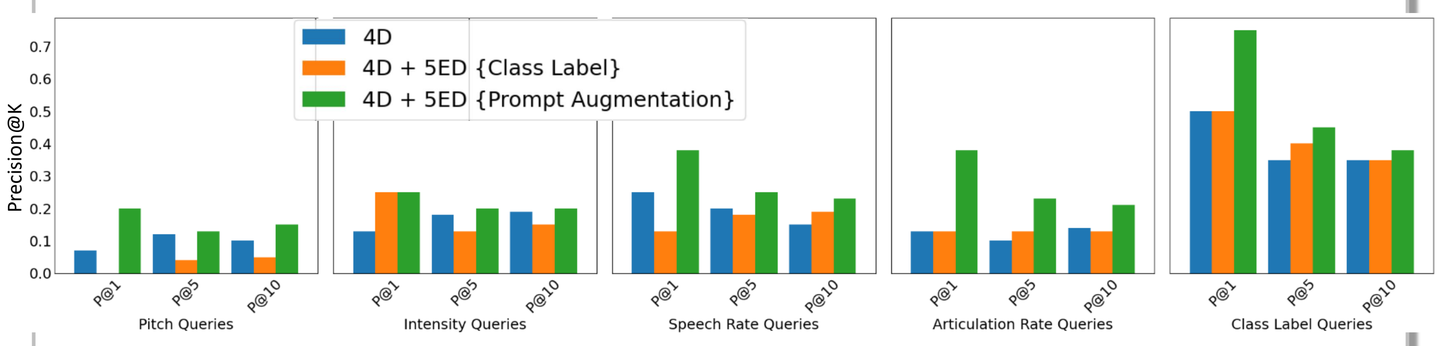

RESULTS – Speech Emotion Recognition

- The performance is shown on RAVDESS dataset. The dataset is not included in the training datasets.

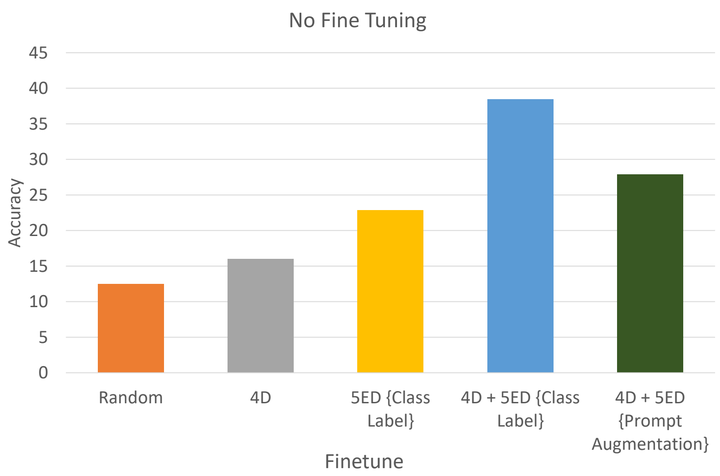

- The model is finetuned on the training subset of RAVDESS and performance shown on testing subset.

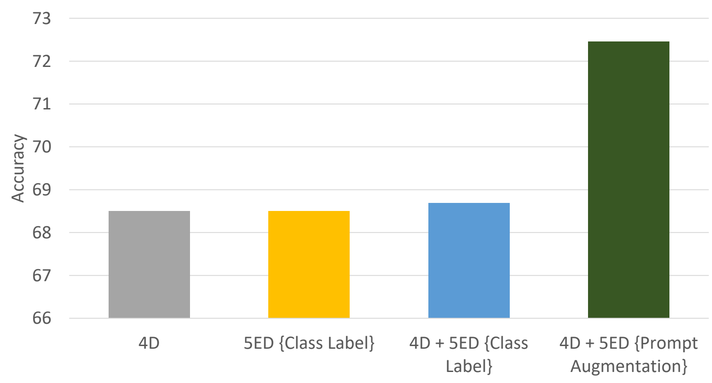

ANALYSIS – which Prompt is better?

- Performance shown when the model is finetuned on RAVDESS dataset using different acoustic prompts paired with the audio.

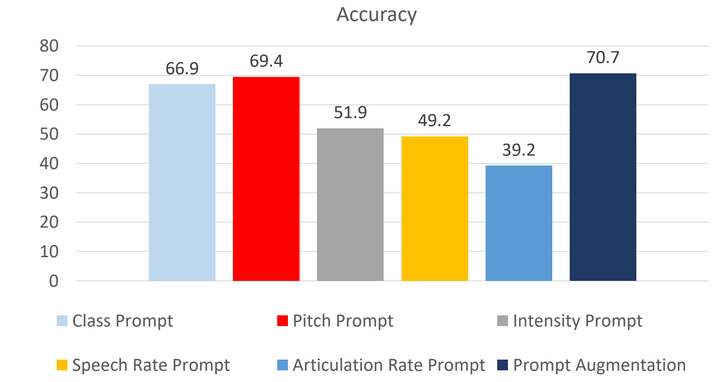

CONCLUSION

- We find that among the acoustic prompts, pitch prompt is the best performing one.
- Emotion Audio Retrieval - acoustic prompt augmentation achieves consistently better Precision@K metric.
- Speech Emotion Recognition - shows performance improvement by 3.8% absolute in RAVDESS.

[1] Scherer, Klaus R. Acoustic concomitants of emotional dimensions: Judging affect from synthesized tone sequences. (1972).
[2] Pavlenko, Aneta. Emotions and multilingualism. Cambridge University Press, 2005.
[3] Benjamin Elizalde, et al "Clap: Learning audio concepts from natural language supervision," arXiv preprint arXiv:2206.04769, 2022.